

Denoising the Observed Series as Preprocessing for Stochastic Optimization: When Filtering Rescues Convergence under Heavy-Tailed Gradient Noise

Dmitry Kortunov
HSE University
Moscow, Russia

Vadim Pastushenko
HSE University
Moscow, Russia

Aleksey Dubovitsky
HSE University
Moscow, Russia

Andrey Shadrin
HSE University
Moscow, Russia

Abstract—We study denoising of the observed time series as an explicit stage of the stochastic optimization pipeline. The observation model is $y_t = s_t + \xi_t$ and a filter F is applied before the optimizer. We give the exact operator and the relevant analytical properties of every noise model, filter and optimizer used, and prove the mechanism behind the empirical finding: under α -stable gradient noise ($\alpha < 2$, infinite variance) constant-step SGD has no finite noise floor and does not converge; a windowed median maps an infinite-variance window to a finite-variance order statistic, which restores the standard $O(\eta L \sigma^2)$ floor, whereas every linear filter (moving average, Kalman) leaves the α -stable law — and the infinite variance — intact. Empirically, on a calibrated tail index $\hat{\alpha} \approx 1.2$ measured on real series, plain SGD on quadratic regression and logistic classification converges in 0/8 seeds (also with the linear and wavelet filters), and in 8/8 with the causal median, lowering the noise floor by 108× and 133×; the effect persists under data-calibrated tails (91×/120×) and on a real, fully-causal 15-minute task ($\approx 78\times$ faster). Negative zones (clean Gaussian noise, long-range-dependent mixtures) are predicted by the same theory and confirmed.

Index Terms—stochastic optimization, heavy-tailed noise, α -stable distributions, median filtering, Kalman filter, wavelet thresholding, long-range dependence, convergence.

I. INTRODUCTION

Convergence theory for stochastic first-order methods rests on a finite second moment of the gradient noise and, usually, on weak temporal dependence. Both fail for learning on financial series, whose returns are heavy-tailed and long-memory [1], [2] and whose stochastic gradients inherit these properties [3]. We treat denoising of the observed series not as visualization but as a computational stage placed between the observations and the optimizer, and ask a precise question: for which (noise, filter, optimizer) triples does preprocessing turn a non-convergent run into a convergent one, and why. The contribution is (i) a self-contained derivation of the operator and the relevant properties of every component, (ii) a proposition isolating the mechanism — a nonlinear order-statistic filter restores a finite effective noise variance that linear filters cannot — and (iii) an empirical confirmation with confidence intervals over seeds, on synthetic, data-calibrated and real causal regimes, including the faithfully reported regimes where filtering does not help.

II. RELATED WORKS

The joint heavy-tailed / long-range-dependent (LRD) noise regime in stochastic approximation was first analyzed asymptotically by Anantharam and Borkar [4] through a limiting ODE comparison; Chandak *et al.* [5] give the first finite-time moment bounds for the same regime. Both analyze the unmodified iteration. Robustness to heavy tails is otherwise obtained by modifying the optimizer: gradient clipping [6], high-probability analyses under unbounded variance [7], [8], and nonlinear SGD [9]. Classical denoising operators — the Kalman filter [10], wavelet soft-thresholding [11]–[13], weighted medians [14] — and the decorrelation of LRD processes by wavelets [15] are classical; denoising as a preprocessing step for forecasting was reported empirically by [16]. Our question is orthogonal to all of the above: we keep the optimizer fixed and quantify, with a mechanism, *where and by how much* preprocessing the observed series rescues convergence.

III. PROBLEM STATEMENT

A. Learning problem

Let $x \in \mathbb{R}^d$ be the model parameters. From the observed scalar series $y_{1:T}$ we form supervised pairs $(u_i, v_i)_{i=1}^n$, where the input $u_i \in \mathbb{R}^p$ is a lagged window and the target v_i is the next value (or its sign for classification). With a model g_x and a loss ℓ (squared for regression, logistic cross-entropy for classification) we minimize the empirical risk

$$f(x) = \frac{1}{n} \sum_{i=1}^n \ell(g_x(u_i), v_i), \quad x^* \in \arg \min_{x \in \mathbb{R}^d} f(x). \quad (1)$$

B. Stochastic step and the place of the filter

A first-order method has access at iteration k to a noisy gradient

$$g_k = \nabla f(x_k) + \xi_k, \quad x_{k+1} = x_k - \eta_k \phi(g_k), \quad (2)$$

with ξ_k a centered noise; we assume neither $\mathbb{E} \|\xi_k\|^2 < \infty$ nor independence across k . The observation model is

$$y_t = s_t + \xi_t, \quad (3)$$

s the signal, ξ the noise. A filter is a map $\hat{s} = F(y_{1:T})$ applied *before* the optimizer; equivalently it acts on the injected

gradient noise. The optimizer then runs on $\hat{s}$. Input of the whole procedure: $y_{1:T}$; output: x_K and convergence/quality metrics. The control is $F_0 = \text{Id}$ (no filtering). Thus F sits strictly between (3) and (2).

IV. NOISE MODELS

All noises are centered; $\sigma > 0$ is a scale.

N1 Gaussian (control): $\xi_t \sim \mathcal{N}(0, \sigma^2)$ i.i.d.: finite variance, independent — the regime classical theory assumes.

N2 long memory (FARIMA): $\xi_t = \sum_{j \geq 0} \psi_j \eta_{t-j}$, $\eta_j \sim \mathcal{N}(0, \sigma^2)$, ψ_j the FARIMA(0, d , 0) weights $\psi_j = \Gamma(j + d)/(\Gamma(d)\Gamma(j + 1)) \sim j^{d-1}/\Gamma(d)$, $d \in (0, \frac{1}{2})$. The variance is finite but the autocovariance decays as $\gamma(\tau) \sim c\tau^{2d-1}$, so $\sum_{\tau} \gamma(\tau) = \infty$: long-range dependence, spectral density $\propto |\omega|^{-2d}$ (1/ f -type).

N3 heavy tails (α -stable): $\xi_t \sim \mathcal{S}_\alpha(\sigma, 0, 0)$ i.i.d., $\alpha \in (1, 2)$, with characteristic function $\mathbb{E}e^{i\theta\xi} = \exp(-\sigma^\alpha|\theta|^\alpha)$. Then $\mathbb{E}|\xi|^p < \infty$ iff $p < \alpha$; in particular $\mathbb{E}\xi = 0$ but $\text{Var}\xi = \infty$, and the tail is $\Pr(|\xi| > z) \sim C_\alpha \sigma^\alpha z^{-\alpha}$.

N4 mixed (heavy + long memory): $\xi_t = \sum_{j \geq 0} \psi_j \zeta_{t-j}$ with ζ_j i.i.d. $\mathcal{S}_\alpha$ and ψ_j as in N2: infinite variance and long memory simultaneously. Auxiliary regimes (regime switch, jump diffusion, Hawkes clusters) are used only for stress tests.

The single property that drives every result below is: *N3/N4 have infinite second moment* ($\alpha < 2$), while *N1/N2 have a finite one*.

V. FILTERS

Let the per-coordinate noise sequence entering (2) be $\xi_{1:T}$. We give each operator and the one property that matters.

A. F_0 Identity

$\hat{s}_t = y_t$. Baseline; the gradient noise is exactly ξ_t .

B. F_1 Moving average (causal, window w)

$$\hat{s}_t = \frac{1}{w} \sum_{i=0}^{w-1} y_{t-i}. \quad (4)$$

Linear and unbiased. Its key property is closure: a finite linear combination of α -stable variables is α -stable. If $\xi_t \sim \mathcal{S}_\alpha(\sigma, 0, 0)$ i.i.d., then by the stability property the filtered noise $\bar{\xi}_t = w^{-1} \sum_{i=0}^{w-1} \xi_{t-i} \sim \mathcal{S}_\alpha(\sigma_w, 0, 0)$ with

$$\sigma_w = \sigma \left(\sum_{i=0}^{w-1} w^{-\alpha} \right)^{1/\alpha} = \sigma w^{1/\alpha-1}. \quad (5)$$

Since $1 < \alpha < 2$, $1/\alpha - 1 \in (-\frac{1}{2}, 0)$: averaging shrinks the scale only as $w^{1/\alpha-1}$, *slower* than the Gaussian $w^{-1/2}$, and crucially the law remains α -stable with $\text{Var}\bar{\xi} = \infty$.

C. F_2 Kalman local level (causal)

State model $s_t = s_{t-1} + w_t$, observation $y_t = s_t + \varepsilon_t$; the steady-state filter is the exponential recursion

$$\hat{s}_t = (1 - K)\hat{s}_{t-1} + Ky_t = K \sum_{j \geq 0} (1 - K)^j y_{t-j}, \quad (6)$$

$K \in (0, 1)$ the steady-state gain. This is an IIR *linear* operator: an infinite geometric mixture of the inputs. By the same stability/closure argument, with i.i.d. α -stable input the output is α -stable with scale $\sigma(K^\alpha \sum_{j \geq 0} (1 - K)^{\alpha j})^{1/\alpha} = \sigma K(1 - (1 - K)^\alpha)^{-1/\alpha}$, again with infinite variance. Optimal for Gaussian s [10]; powerless against heavy tails.

D. F_3 Wavelet soft-threshold

With Daubechies analysis, soft-threshold the detail coefficients, inverse-transform:

$$\hat{c} = \text{sgn}(c)(|c| - \lambda)_+, \quad \lambda = \hat{\sigma} \sqrt{2 \ln N}, \quad (7)$$

$\hat{\sigma} = \text{MAD}(c^{(\text{finest})})/0.6745$ [11], [12]. The universal factor $\sqrt{2 \ln N}$ is exactly the maximum of N i.i.d. *Gaussian* deviations; it is the wrong order for α -stable extrema, whose maximum over N samples grows as $N^{1/\alpha} \gg \sqrt{\ln N}$. Moreover only detail bands are thresholded: a heavy-tailed spike projects a non-vanishing mass onto the untouched approximation (low-pass) band and survives reconstruction. Hence the residual after F_3 stays heavy-tailed (infinite variance) under N3.

E. F_4 Causal median (window w)

$$\hat{s}_t = \text{med}\{y_{t-w+1}, \dots, y_t\}. \quad (8)$$

This is the only *nonlinear* filter we keep. It is an order statistic of the window and, as shown next, maps an infinite-variance window to a finite-variance output. We restrict the study to F_0 - F_4 ; only causal filters (F_0, F_1, F_2, F_4) are used on real data (no look-ahead).

VI. OPTIMIZERS

Update (2) with ϕ : **SGD** $\phi(g) = g$ (constant step η); **Clipped-SGD** $\phi(g) = \min(1, \tau/\|g\|)g$; **Normalized-SGD** $\phi(g) = g/\|g\|$; **Adam** [17] with first/second-moment EMAs; **AdamW** [18] (decoupled L_2). The relevant fact for constant-step SGD on an L -smooth, μ -strongly convex f is the classical noise-floor identity: if $\mathbb{E}\|\xi_k\|^2 \leq \sigma^2$ then

$$\mathbb{E}\|\nabla f(x_K)\|^2 \xrightarrow{K \rightarrow \infty} \Theta(\eta L \sigma^2), \quad (9)$$

a *finite* plateau proportional to σ^2 [19]. Clipping/normalization replace g by a bounded surrogate, so σ^2 in (9) is finite by construction [6], [7], [9]; this is the optimizer-side route. Our object is the orthogonal, preprocessing-side route through F .

VII. CONVERGENCE METRICS

For each (model, noise, optimizer, filter) cell we run 8 fixed seeds and report, with 95% bootstrap CIs over seeds: (i) *binary convergence* — the fraction of seeds reaching a $10^2\times$ reduction of $\|\nabla f\|^2$ within the horizon; (ii) $T(\varepsilon) = \min\{k : \|\nabla f(x_k)\|^2 \leq \varepsilon\}$; (iii) the *noise floor* — $p_{10}/p_{50}/p_{90}$ of the trailing $\|\nabla f\|^2$; (iv) the divergence slope — the least-squares slope of $\log\|g_k\|^2$ vs. $\log k$ on the last 30% (> 0 : diverging, ≈ 0 : stalled at a floor, < 0 : still converging); (v) the convergence AUC, $\log_{10}\|\nabla f\|^2$; (vi) held-out error (no train/test mixing); (vii) wall-clock including the filter cost.

VIII. THEORETICAL RESULTS

We isolate why a windowed median rescues SGD while linear/wavelet filters do not. Take N3: $\xi_t \sim \mathcal{S}_\alpha(\sigma, 0, 0)$ i.i.d., symmetric, with density h , $h(0) > 0$.

Lemma 1 (Order statistic has finite variance). *Let $\xi_{(1)}, \dots, \xi_{(w)}$ be the order statistics of w i.i.d. draws of a symmetric law with continuous density h , $h(0) > 0$, and m_w the sample median (w odd). Then $\mathbb{E}m_w = 0$ and, as $w \rightarrow \infty$,*

$$\text{Var } m_w = \frac{1}{4w h(0)^2} + o(1/w) < \infty, \quad (10)$$

for every $\alpha \in (0, 2]$, including the infinite-variance α -stable case.

Proof. Symmetry gives $\mathbb{E}m_w = 0$. The classical sample-median CLT $\sqrt{w}(m_w - \theta) \Rightarrow \mathcal{N}(0, (4h(\theta)^2)^{-1})$ holds whenever h is continuous and positive at the population median $\theta = 0$; it uses only the local behavior of h at 0 and no moment of ξ . The α -stable density is bounded, continuous and strictly positive at 0, so (10) follows; finiteness is independent of the (infinite) variance of ξ . $\square$

Proposition 1 (Median preprocessing restores the SGD floor). *Under N3 with the causal median (8) of fixed odd width w , the effective gradient noise $\tilde{\xi}_t = m_w(\xi_{t-w+1:t})$ satisfies $\sigma_{\text{med}}^2 := \mathbb{E}\tilde{\xi}_t^2 \leq \frac{1}{4h(0)^2} (1 + o(1)) < \infty$. Consequently constant-step SGD on an L -smooth μ -strongly convex f admits the finite floor (9) with $\sigma^2 = \sigma_{\text{med}}^2$ and reaches it at the standard rate; in particular the $10^2\times$ reduction event has probability $\rightarrow 1$. For any linear filter $F \in \{F_1, F_2\}$ the filtered noise is again α -stable (Sec. V), so $\mathbb{E}\tilde{\xi}_t^2 = \infty$, (9) is vacuous, and SGD has no finite plateau. For F_3 the residual retains an α -stable component (untouched approximation band; Gaussian-calibrated threshold), hence also $\mathbb{E}\tilde{\xi}_t^2 = \infty$.*

Proof. Lemma 1 bounds the per-window median variance uniformly in α ; for a fixed window the marginal second moment is $\leq (4h(0)^2)^{-1} (1 + o(1))$. With a finite effective variance the constant-step SGD descent lemma yields, after telescoping, $\frac{1}{K} \sum_{k < K} \mathbb{E} \|\nabla f(x_k)\|^2 \leq \frac{2(f(x_0) - f^*)}{\eta K} + \eta L \sigma_{\text{med}}^2$, i.e. the floor (9); strong convexity turns the average into a last-iterate plateau. The linear-filter claims are the closure identities of Sec. V ((5) and its Kalman analogue): stability of the α -stable family is preserved under finite or geometric linear combinations, so the variance stays infinite. The F_3 claim

TABLE I
N3 ($\alpha=1.2$), SGD, 8 SEEDS. BINARY CONVERGENCE AND MEDIAN NOISE FLOOR. $F_1/F_2/F_3$ BEHAVE LIKE F_0 (ALL 0/8).

model	bin. F_0	bin. F_4	floor F_0	floor F_4	ratio
quadratic	0/8	8/8	0.540	0.0050	108 $\times$
logistic	0/8	8/8	0.0658	0.00050	133 $\times$
AR(5)	8/8	8/8	0.0421	0.0021	20 $\times$

follows because soft-thresholding the detail bands leaves the low-pass projection of a heavy-tailed atom, whose α -stable tail is unchanged. $\square$

Remark 1 (Gaussian control — a predicted negative). *Under N1 the same Lemma gives $\text{Var } m_w = (4wh(0)^2)^{-1} = \pi\sigma^2/(2w)$ while the mean has σ^2/w ; the asymptotic relative efficiency of the median is $2/\pi \approx 0.637$. Filtering clean Gaussian noise therefore raises the floor by $\pi/2 \approx 1.57$: preprocessing must not be applied when $\alpha \approx 2$.*

Remark 2 (Mixed N4 — a predicted negative). *Lemma 1 needs (near-)independent window samples. Under N4 the window is strongly dependent ($\sum_\tau \gamma(\tau) = \infty$) and heavy-tailed jumps cluster; the sample-median CLT variance is inflated and (10) no longer bounds the floor. The median does not rescue the mixed regime.*

Remark 3 (Adam under long memory N2). *N2 has finite variance, so (9) already holds; the issue is the bias of Adam's second-moment EMA under the persistent $1/f$ autocorrelation, which violates the near-independence used in Adam's bounds [20]. Wavelet bases approximately decorrelate LRD processes across scales [15], returning the filtered noise to a near-white regime; hence F_3 (not F_4) accelerates Adam/AdamW here.*

IX. EXPERIMENTS

A. Setup

Models: quadratic regression ($d=20$), logistic classification ($d=15$), AR(5). Noise scales fixed across all filters/optimizers; 8 seeds; pre-registered, no per-seed/window/metric selection. Tail indices are estimated on genuinely real series (McCulloch $\hat{\alpha}$, DFA $\hat{H}$ on $|r|$): SPY $\hat{\alpha}=1.05$, $\hat{H}=0.86$ (the -11.6% COVID drop of 2020-03-16 is present), basket median $\hat{\alpha}=1.21$, $\hat{d}=0.29$.

B. Heavy-tailed SGD: only the median rescues

Table I: under N3 ($\alpha=1.2$) plain SGD and SGD with every linear/wavelet filter converge in 0/8 seeds; the causal median converges in 8/8 and lowers the floor by 108 $\times$ (quadratic) and 133 $\times$ (logistic); for AR all cells converge but the median floor is 20 $\times$ lower. On the Gaussian control N1 the median floor ratio is 0.76–0.95 (< 1 : a mild loss), exactly the $\pi/2$ efficiency penalty. Fig. 1 shows the floor and the median $\|\nabla f\|^2$ curves.

Предфильтрация спасает сходимость SGD при тяжёлохвостовом градиентном шуме (8 сидов)




Fig. 1. Heavy-tailed N3, SGD, 8 seeds. (a) noise floor F_0 – F_4 per model (log scale, p_{10} – p_{90}); (b) median $\|\nabla f(x_k)\|^2$: F_0 / F_2 stall high, the causal median F_4 drops 1–2 orders. Linear and wavelet filters track F_0 , as Proposition 1 predicts.

TABLE II
APPLIED CAUSAL AR(5) TRAINING, ADAM, REAL DOMAINS. T : MEDIAN ITERATIONS TO A $10^2 \times$ GRADIENT DROP (CAP 4000).

domain	$T(F_0)$	conv. F_0	best	speed-up
fin. 15-min	4000	7/16	$F_2 \approx 52$	78 $\times$
fin. daily	554	13/16	$F_2 \approx 57$	10 $\times$
sensor (ETT)	43	16/16	F_1 20	2.1 $\times^\dagger$
macro (FRED)	20	16/16	F_0	—

† faster but held-out MSE worsens (honest negative).

C. Data-calibrated tails

At the empirically measured $\hat{\alpha}=1.21$ the rescue persists: F_4 lowers the SGD floor by $91\times$ (quadratic), $120\times$ (logistic), $17\times$ (AR). At the calibrated mixed tail ($\hat{d}=0.29$, $\hat{\alpha}=1.21$) the gain is $\approx 1\times$ — the predicted N4 negative (Remark 3).

D. Adaptive optimizers

On long-memory N2 (quadratic) wavelet F_3 accelerates Adam and AdamW identically by $11.2\times$ ($T(\varepsilon)$: $562 \rightarrow 50$; F_2 $3.1\times$); on the mixed/logistic cells no filter beats F_0 — the predicted Remark 4 picture.

E. Applied, fully causal

Table II: on real 15-minute returns Adam without a filter converges in only $\approx 7/16$ runs (median $T=4000$, the cap), and in $16/16$ with the causal Kalman filter ($T \approx 52$, $\approx 78\times$), at unchanged held-out MSE (0.528 vs. 0.522); features use only the past and the held-out target is the raw future, identical across filters. Daily returns: $\approx 10\times$. Faithfully reported negatives: the sensor series speeds up but loses held-out accuracy (oversmoothing); on near-white raw daily returns and on the offline macro fallback F_0 is best.

F. Recoverable-signal regimes

When a smooth signal exists under the real noise the gain is large and unambiguous: with real returns as noise plus a

known signal, wavelet F_3 recovers it at $8.1\times$ lower held-out MSE; in a microstructure simulation the median lowers the integrated-variance error by $\approx 58\times$.

X. CONCLUSION

The effect is mechanistic, not incidental: α -stable gradient noise removes the finite second moment that constant-step SGD needs, linear and wavelet filters preserve the α -stable law, and only a nonlinear order statistic — the windowed median — restores a finite effective variance and hence the standard noise floor. This predicts, and the experiments confirm with CIs over seeds, both the positive regime (heavy tails: $0/8 \rightarrow 8/8$, floor 108 – $133\times$ lower; real 15-min: $\approx 78\times$) and the negative regimes (clean Gaussian: $\pi/2$ penalty; long-memory mixtures: no gain). The practical rule follows directly: estimate $\hat{\alpha}, \hat{H}$ on the training span; apply the causal median only when tails are heavy and a recoverable signal exists; never validate with look-ahead.

REFERENCES

- [1] R. Cont, “Empirical properties of asset returns: Stylized facts and statistical issues,” *Quantitative Finance*, vol. 1, no. 2, pp. 223–236, 2001.
- [2] B. Mandelbrot, “The variation of certain speculative prices,” *The Journal of Business*, vol. 36, no. 4, pp. 394–419, 1963.
- [3] U. Şimşekli, L. Sagun, and M. Gürbüzbalaban, “A tail-index analysis of stochastic gradient noise in deep neural networks,” in *Proceedings of the 36th International Conference on Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 97. PMLR, 2019, pp. 5827–5837. [Online]. Available: <https://arxiv.org/abs/1901.06053>
- [4] V. Anantharam and V. S. Borkar, “Stochastic approximation with long range dependent and heavy tailed noise,” *Queueing Systems*, vol. 71, no. 1–2, pp. 221–242, 2012.
- [5] S. Chandak, A. K. Yadav, A. Özgür, and N. Bambos, “Heavy-tailed and long-range dependent noise in stochastic approximation: A finite-time analysis,” Mar. 2026, submitted to IEEE Transactions on Automatic Control. [Online]. Available: <https://arxiv.org/abs/2603.19648>
- [6] E. Gorbunov, M. Danilova, and A. Gasnikov, “Stochastic optimization with heavy-tailed noise via accelerated gradient clipping,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, 2020, pp. 15 042–15 053. [Online]. Available: <https://arxiv.org/abs/2005.10785>

- [7] A. Cutkosky and H. Mehta, “High-probability bounds for non-convex stochastic optimization with heavy tails,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2021. [Online]. Available: <https://arxiv.org/abs/2106.14343>
- [8] A. Sadiev, M. Danilova, E. Gorbunov, S. Horváth, G. Gidel, P. Dvurechensky, A. Gasnikov, and P. Richtárik, “High-probability bounds for stochastic optimization and variational inequalities: The case of unbounded variance,” in *Proceedings of the 40th International Conference on Machine Learning (ICML)*, 2023. [Online]. Available: <https://arxiv.org/abs/2302.00999>
- [9] A. Armacki, S. Yu, P. Sharma, G. Joshi, D. Bajović, D. Jakovetić, and S. Kar, “Nonlinear stochastic gradient descent and heavy-tailed noise: A unified framework and high-probability guarantees,” 2024. [Online]. Available: <https://arxiv.org/abs/2410.13954>
- [10] R. E. Kalman, “A new approach to linear filtering and prediction problems,” *Journal of Basic Engineering*, vol. 82, no. 1, pp. 35–45, 1960.
- [11] D. L. Donoho and I. M. Johnstone, “Ideal spatial adaptation by wavelet shrinkage,” *Biometrika*, vol. 81, no. 3, pp. 425–455, 1994.
- [12] D. L. Donoho, “De-noising by soft-thresholding,” *IEEE Transactions on Information Theory*, vol. 41, no. 3, pp. 613–627, 1995.
- [13] S. Mallat, *A Wavelet Tour of Signal Processing*, 2nd ed. San Diego: Academic Press, 1999.
- [14] L. Yin, R. Yang, M. Gabbouj, and Y. Neuvo, “Weighted median filters: A tutorial,” *IEEE Transactions on Circuits and Systems II: Analog and Digital Signal Processing*, vol. 43, no. 3, pp. 157–192, 1996.
- [15] D. B. Percival and A. T. Walden, “Wavelet methods for time series analysis,” *Cambridge Series in Statistical and Probabilistic Mathematics*, 2000.
- [16] Q. Tang, T. Fan, R. Shi, J. Huang, and Y. Ma, “Prediction of financial time series using LSTM and data denoising methods,” 2021. [Online]. Available: <https://arxiv.org/abs/2103.03505>
- [17] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” in *Proceedings of the 3rd International Conference on Learning Representations (ICLR)*, 2015. [Online]. Available: <https://arxiv.org/abs/1412.6980>
- [18] I. Loshchilov and F. Hutter, “Decoupled weight decay regularization,” in *Proceedings of the 7th International Conference on Learning Representations (ICLR)*, 2019. [Online]. Available: <https://arxiv.org/abs/1711.05101>
- [19] H. Robbins and S. Monro, “A stochastic approximation method,” *The Annals of Mathematical Statistics*, vol. 22, no. 3, pp. 400–407, 1951.
- [20] A. Défossez, L. Bottou, F. Bach, and N. Usunier, “A simple convergence proof of adam and adagrad,” *Transactions on Machine Learning Research*, 2022. [Online]. Available: <https://arxiv.org/abs/2003.02395>